

Zoom Phone Transcript Pipeline

Automated delivery of Zoom Phone call transcripts into SharePoint

Sensible Care Pty Ltd

1. What this system does

Zoom Phone records and transcribes calls, but keeps those transcripts only inside Zoom, where they are awkward to search, impossible to retain under policy, and invisible to the rest of Microsoft 365. This pipeline copies every transcript into a SharePoint document library in Sensible Care's own tenant, filed by the team that took the call and the week it happened.

Three things follow from that. Transcripts become searchable alongside everything else in SharePoint; Microsoft Purview retention labels can be applied to them like any other document; and no third party ever holds the data, because it moves directly from Zoom to Sensible Care's tenant.

Property	Value
Source	Zoom Phone recording transcripts, via the Zoom API
Destination	SharePoint document library <i>Call Transcripts</i>
Filing	<code>{Team}/{ISO_year}-W{week}/</code>
Time zone	Australia/Melbourne — weeks are local, not UTC
Format	Plain text (<code>.txt</code>) with a metadata header
Volume	Roughly 640 transcripts per day
History	Transcription was enabled around 21 July 2025; nothing exists before that
Hosting	Azure Container Apps, region <code>australiaeast</code>

2. The three moving parts

The same delivery logic runs in three different places. They differ only in what makes them start and which dates they cover.

The webhook — real time

An always-on web service. Zoom calls it the moment a transcript finishes processing, and the transcript is usually in SharePoint within seconds. This is the primary path; almost everything arrives this way.

The nightly job — the safety net

Runs once a day and re-checks the last two days. Anything the webhook missed — because Zoom's delivery failed, or the service was restarting — is picked up here. It then seals the previous week's manifests and reconciles what Zoom holds against what SharePoint holds.

The backfill job — history

Walks backwards through completed months, one calendar month per pass, delivering everything recorded before the pipeline existed. It is designed to be killed and restarted at any moment. When the historical archive is complete this job

is switched off and never runs again.

3. How a transcript travels

1. **Zoom signals.** A `phone.recording_transcript_completed` event is posted to the webhook, or a scheduled job queries `GET /v2/phone/recordings` for a date range.
2. **The sender is verified.** Webhook requests carry an HMAC-SHA256 signature over `v0:{timestamp}:{body}`. A bad signature, or a timestamp too far from now, is rejected as a possible replay.
3. **The owner is resolved to a team.** The recording names a Zoom user or a call queue. That is turned into an email address, and the email is looked up in the team map.
4. **Prior delivery is checked.** The destination week is listed in SharePoint. If this recording's id is already there, it is skipped.
5. **The transcript is fetched and rendered.** Zoom returns JSON; the pipeline converts it to readable text and prepends a metadata header.
6. **It is uploaded** to `{Team}/{ISO_year}-w{week}/`, creating the folder if needed.
7. **A manifest row is recorded** for the weekly summary.

Transcript text is never written to local disk. It is streamed from Zoom, rendered in memory, and sent to SharePoint. Nothing remains on the machine running the pipeline.

4. Where files land

```
Call Transcripts/  
  
  Care Partner Team/  
  
    2026-W39/  
  
      143301_inbound_+614000000000_to_200_d8d963272aa1.txt  
  
      _manifest_2026-W39_Care Partner Team.csv  
  
  Rostering/  
  
    2026-W39/  
  
    _Unassigned/  
  
      2026-W39/  
  
      _config/  
  
      teams.csv
```

Filenames are `HHMMSS_direction_caller_to_callee_shortid.txt`, in Melbourne local time. Leading with the time makes a folder sort chronologically, and the twelve-character recording id on the end makes collisions impossible.

Why ISO weeks and not months. The folder name pairs the ISO year with the ISO week, never the calendar

year. 1 January 2027 falls in ISO week 53 of 2026; writing 2027-W53 would create a folder that should not exist. The code enforces the pairing.

5. Weekly manifests

Each team gets one CSV per week listing every transcript delivered: recording id, time, direction, duration, team, filename, and when it was uploaded.

A manifest is written only once the week has closed and a grace period has passed, so that late arrivals are included. Once written it is never rewritten. Anything that arrives afterwards goes into a numbered supplement file beside it, containing only the genuinely new rows.

6. Azure infrastructure

Resource	Name	Purpose
Resource group	rg-zoom-transcripts	Holds everything
Container Apps environment	cae-zoom-transcripts	Shared network and logging
Container app	app-transcripts-webhook	Always-on webhook receiver
Job (scheduled)	job-transcripts-nightly	Daily safety net, seal, reconcile
Job (scheduled)	job-transcripts-backfill	Historical months
Job (manual)	job-transcripts-manual	Ad-hoc operator runs
Container registry	Azure Container Registry	Stores the built image
File share	Azure Files, mounted at /data	State that must survive restarts
Logs	Log Analytics	Console output from every run

Containers are destroyed when they finish. Everything that must outlive a run is written to the Azure Files share, which persists independently of them.

7. Schedules

Azure runs cron in **UTC and does not observe daylight saving**. The UTC column is what is configured; the Melbourne column is the same instant expressed locally, and shifts when Melbourne changes clocks.

Job	UTC	Melbourne
Nightly	16:00 daily	02:00 next day (03:00 in daylight time)
Backfill	11:00 daily, up to 22 h	21:00 (22:00 in daylight time)
Webhook	Always running; no schedule	